

**leti****list**

## **Flexible and Performing Kernels Dynamically Generated with deGoal for Embedded Systems**

**Damien Couroussé<sup>(1)</sup>**

**Henri-Pierre Charles<sup>(1)</sup>**

**Yves Lhuillier<sup>(2)</sup>**

`author.name@cea.fr`

CEA-LIST, LaSTRE/Grenoble (1) and LCE (2) laboratories

December 1st, 2011

# Outline

## 1 Overview of deGoal

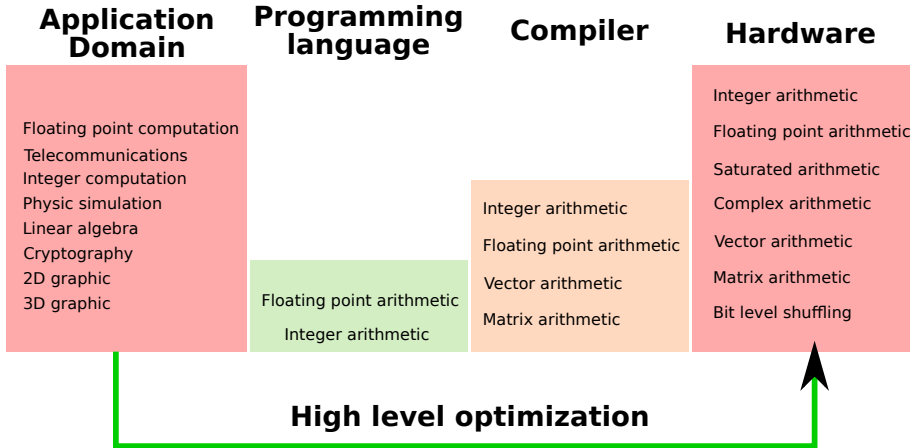
- motivation
- a high level assembly language
- the development workflow

## 2 Two application cases

- Matrix multiplication
- A memory allocator for embedded systems

## 3 Conclusion

# the semantic bottleneck



# motivation for developing deGoal

- 1 the semantic bottleneck of C language
- 2 limitations of static compilers
  - vector instructions: limited use
  - target's ISA is partially covered. eg: psadd (Cell), fpx\_MAC (FPx on STxP70), etc.
  - **no code optimizations from execution context** (data values, memory addresses, etc.)
- 3 Claim: C source code for high performance is *not* portable
  - architecture-specific intrinsics
  - compiler-dependent pragmas
  - the final binary code depends on the compiler technology: gcc, icc, open64, LLVM, etc.

# deGoal: Dynamic Execution of Generated Optimized kernel from Algorithmic Level

- code generation depends on the execution context
  - constant propagation from runtime data (values and memory addr.)
  - instruction selection
  - removal of dead code
- support of application-specific instructions
- a common language for all architectures
- suitable for many-core embedded systems
  - code generators have a very small footprint
  - very fast code generation (10 to 100 cycles / instruction)
  - supports heterogeneous architectures
  - control over code generation (not done behind developer's back)

## Disclaimer

We focus at the processor's level (i.e. PEs in P2012).

# a high level assembly language

## Characteristics:

Architecture independent

Multiple arithmetics integer, floating point, complex, saturated

Register allocation virtual vector registers of multiple width,  
automatic sizing of vectors (WIP)

Dynamically parametrized mix C expressions with generated code

Multimedia operators SAD, DFT, memory access patterns  
operators are automatically adjusted to target capabilities  
(e.g. availability of the psadd instruction)

No Intermediate Representation for very fast code generation

# the development workflow


**Compilette** the runtime code generator,  
kernel-specific  
designed to be very fast

**Kernel** the code section to optimize,  
generated at runtime

# the development workflow

**Compilette** the runtime code generator,  
kernel-specific  
designed to be very fast

**Kernel** the code section to optimize,  
generated at runtime

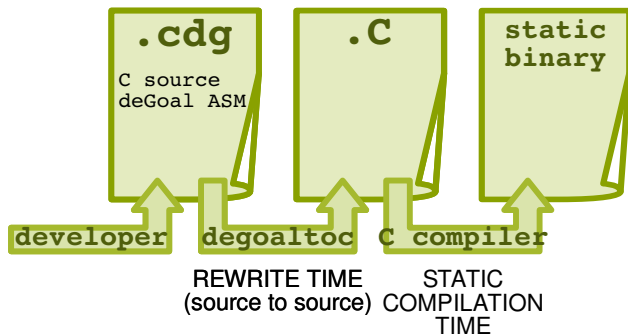




# the development workflow

**Compilette** the runtime code generator,  
kernel-specific  
designed to be very fast

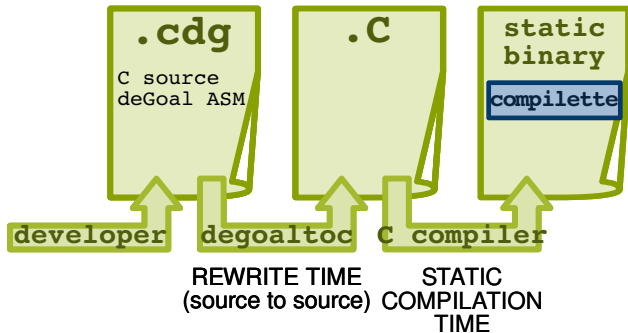
**Kernel** the code section to optimize,  
generated at runtime



# the development workflow

**Compilette** the runtime code generator,  
kernel-specific  
designed to be very fast

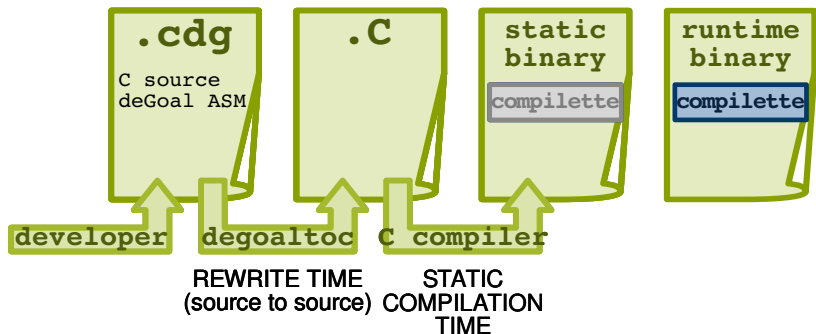
**Kernel** the code section to optimize,  
generated at runtime



# the development workflow

**Compilette** the runtime code generator,  
kernel-specific  
designed to be very fast


**Kernel** the code section to optimize,  
generated at runtime



# the development workflow

**Compilette** the runtime code generator,  
kernel-specific  
designed to be very fast

**Kernel** the code section to optimize,  
generated at runtime



# Outline

## 1 Overview of deGoal

- motivation
- a high level assembly language
- the development workflow

## 2 Two application cases

- Matrix multiplication
- A memory allocator for embedded systems


## 3 Conclusion

# naive implementation

$$[C] = [A] \times [B] \quad (1)$$

$$\forall x \in \{0, \dots, n-1\}, \forall y \in \{0, \dots, q-1\}, c_{xy} = \sum_{i=0}^{p-1} a_{iy} b_{xi} \quad (2)$$

```
clear(C)
for (y=0; y < n; y++) {
  for (x=0; x < q; x++) {
    for (i=0; i < p; i++) {
      C[x,y] =
        C[x,y]+A[i,y]*B[x,i]
    } } }
```




# experiment #1 with deGoal

**Objective:** simple code generation, constant propagation only

- propagation of data constants into the binary code
- loop unrolling

```
// generate kernel
mac = gen_mac(&A, &B, &C);

clear(C)
for (y=0; y < n; y++){
  for (i=0; i < p; i+=VSIZE){
    mac(y,i); // execute kernel
  } }
```




# experiment #1 with deGoal

**Objective:** simple code generation, constant propagation only

- propagation of data constants into the binary code
- loop unrolling

```
// generate kernel
mac = gen_mac(&A, &B, &C);

clear(C)
for (y=0; y < n; y++){
  for (i=0; i < p; i+=VSIZE){
    mac(y,i); // execute kernel
  } }
```






# experiment #1 with deGoal

**Objective:** simple code generation, constant propagation only

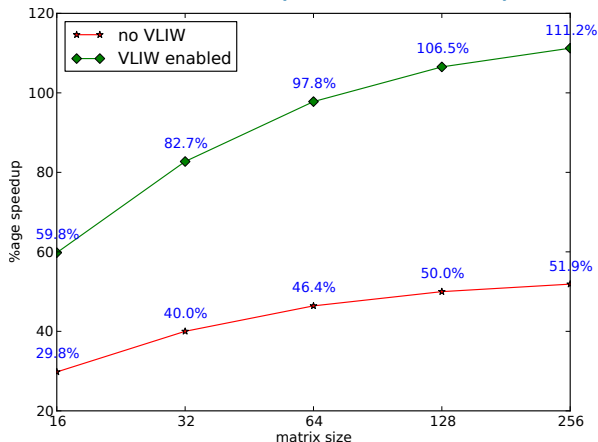
- propagation of data constants into the binary code
- loop unrolling

```
// generate kernel
mac = gen_mac(&A, &B, &C);

clear(C)
for (y=0; y < n; y++){
  for (i=0; i < p; i+=VSIZE){
    mac(y,i); // execute kernel
  } }
```




## experiment #1: results (integer data)



- STxP70-v4 core, simulation with `sxrun` in **CAS** mode
- `-O3, -funroll-loops -Munroll -Mconfig=mult:yes,`
- VLIW enabled with `-Mconfig=vliw:dualcoreALU`

## experiment #1: results (float data)



- STxP70-v4 core, simulation with `sxrun` in **ISS mode**
- `-O3, -funroll-loops -Munroll -Mconfig=mult:yes,`
- VLIW enabled with `-Mconfig=vliw:dualcoreALU`

## experiment #2 with deGoal


### Objective:

- same as #1
- code generation depending on the **data values in A**

```
clear(C)

// generate the generic section
mac=gen_mac(prog_buf,&A,&B,&C);

// process matrix multiplication
for (y=0; y < n; y++){
  for (i=0; i < p; i+=VSIZE){
    // do data specialization
    mac=gensp_mac(prog_buf,y,i);
    if (NULL != mac)
      mac(y,i);
  } }
```



## experiment #2 with deGoal


### Objective:

- same as #1
- code generation depending on the **data values in A**

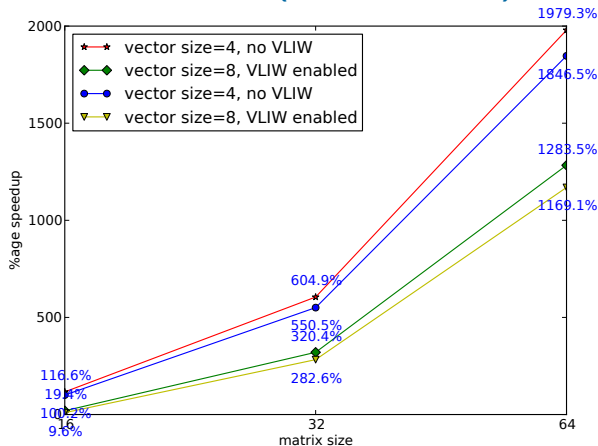
```
clear(C)

// generate the generic section
mac=gen_mac(prog_buf,&A,&B,&C);

// process matrix multiplication
for (y=0; y < n; y++){
  for (i=0; i < p; i+=VSIZE){
    // do data specialization
    mac=gensp_mac(prog_buf,y,i);
    if (NULL != mac)
      mac(y,i);
  } }
```



## experiment #2: results (integer data)



- STxP70-v4 core, simulation with `sxrun` in **CAS** mode
- same compilation options as experiment #1
- results displayed = best case: matrix A is the identity matrix

# Outline

## 1 Overview of deGoal

- motivation
- a high level assembly language
- the development workflow

## 2 Two application cases

- Matrix multiplication
- A memory allocator for embedded systems

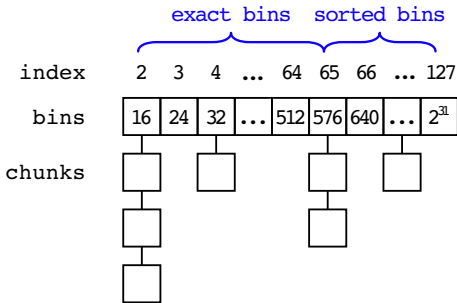
## 3 Conclusion

# an overview of dlmalloc

reference: dlmalloc, main  
implementation in GNU/Linux  
systems

- chunk: reference to a fragment of free memory
- bins: array of linked lists
- the bin\_hash function associates a fragment size to a index (binary tree)

data structure used in dlmalloc  
(stolen from<sup>a</sup>):



---

<sup>a</sup>D. Lea: A memory allocator,  
<http://g.oswego.edu/dl/html/malloc.html>, 2000.

malloc processing:

- 1 index = bin\_hash(size)
- 2 return best\_fit(bins[index], size)



# requirements

... for a memory allocator targeting many-core embedded systems:

- ability to deal with multiple physical memories (TCMs, intra-cluster shared mem, inter-cluster memory ...)
- a multi-instanciable allocator, each instance being parameterized according to its execution context
- a configurable `bin_hash` function; the hash table depends on:
  - the architecture / memory organization
  - the application profile


for further details, follow Yves's presentation

"A Portable Runtime for different P2012 hardware flavors"


16:15 – room 225

# performance results

**Execution time of the hash function (nb. cycles):**



**Speedup factor:**  
(min/average/max)  
over number of successive  
code generations



# Conclusion

Current results for P2012:

- ported the STxP70 v3 & v4 cores
- FPx and VLIW support (STxP70-4)

Other architectures currently supported by deGoal:

- ARM thumb (T1-4), 32 bits (ARM32\_1-3), NEON SIMD
- MIPS (host processor of CEA-LETI GENEPI architecture)
- CUDA PTX assembly language
- Other NDA-based ABI architectures (RISC/VLIW)

# leti

LABORATOIRE D'ÉLECTRONIQUE  
ET DE TECHNOLOGIES  
DE L'INFORMATION

# list

LABORATOIRE D'INTÉGRATION  
DES SYSTÈMES  
ET DES TECHNOLOGIES

## Thanks! Questions?



énergie atomique • énergies alternatives

